


RAG (Retrieval Augmented Generation)


1. The Problem:



You cannot simply paste 500 pages of documentation into a prompt.

- **Context Window Limits:** Models have limits (e.g., 8k, 128k tokens). A 500-page PDF is ~250k tokens. It won't fit.
- **Cost:** You pay per token. Sending a book for every "Hello" is prohibitively expensive.
- **"Lost in the Middle":** LLMs are great at remembering the beginning and end of a prompt but often hallucinate or forget details buried in the middle of massive text blocks.
- **Solution: RAG.** Only send the 3-5 most relevant paragraphs (chunks) needed to answer the specific question.

2. The RAG "ETL" Pipeline (Extract, Transform, Load)



Before you can chat, you must prepare your data. This is a one-time "ingestion" process.

- **Document Loaders (Extract):** Spring AI provides specific readers to pull text from raw files.
 - **PDF:** `PagePdfDocumentReader` (uses Apache PdfBox).
 - **Text/Markdown:** `TextReader`.
 - **JSON:** `JsonReader` (great for structured data).
- **Text Splitters (Transform):** You must break documents into "bite-sized" pieces (chunks) for the AI.
 - **TokenTextSplitter (The Standard):** Splits by token count (e.g., 800 tokens). Safe for LLM context windows. Preserves context by keeping "edges" of paragraphs intact.
 - **Semantic Chunking:** Advanced. Uses an LLM to decide where a "topic" ends. Slower but higher quality.
 - Use **JChunk** or similar libraries to support different types of chunking: <https://docs.jchunk.io/docs/category/chunkers>
- **Embedding & Storing (Load):**

- Convert chunks to vectors.
- Save to `VectorStore` (Postgres/PGVector).

3. Configuring Retrieval: Accuracy vs. Recall

When the user asks a question, how do we find the "right" chunks? You tune the `SearchRequest` inside the advisor.

- `topK` (**Default: 4**): How many chunks to retrieve.
 - *Low (1-2)*: Cheaper, focused, but might miss context.
 - *High (10+)*: Comprehensive, but risks "noise" and higher cost.
- `similarityThreshold` (**0.0 to 1.0**): The "Quality Gate".
 - *0.7 - 0.8*: High relevance. "Only show me exact matches."
 - *0.5*: Loose relevance. "Show me anything vaguely related."

4. Hands-on: Building a "Spring Boot Q&A Bot"

```
<dependency>
  <groupId>org.springframework.ai</groupId>
  <artifactId>spring-ai-pdf-document-reader</artifactId>
</dependency>
```

Step 1: The Ingestion Service (Run once on startup)

This loads the Spring Boot Reference PDF, splits it, and saves it to the Vector DB.

```
@Component
public class IngestionService {

  private final VectorStore vectorStore;
  @Value("classpath:spring-boot-reference.pdf")
```

```

private Resource pdfResource;

public IngestionService(VectorStore vectorStore) {
    this.vectorStore = vectorStore;
}

public void ingest() {
    // 1. Load: Read PDF
    var pdfReader = new PagePdfDocumentReader(pdfResource);
    List<Document> rawDocs = pdfReader.get();

    // 2. Transform: Split into 800-token chunks
    var splitter = new TokenTextSplitter();
    List<Document> chunks = splitter.apply(rawDocs);

    // 3. Load: Store in PGVector
    vectorStore.add(chunks);
    System.out.println("Ingested " + chunks.size() + " chunks!");
}
}

```

Step 2: The ChatBot (The RAG Endpoint)

We configure the QuestionAnswerAdvisor to use the "Stuff" strategy with strict filtering.

```

@RestController
public class DocChatController {

    private final ChatClient chatClient;
    private final VectorStore vectorStore;

    public DocChatController(ChatClient.Builder builder, VectorStore vectorStor
e) {
        this.vectorStore = vectorStore;
    }
}

```

```
    this.chatClient = builder.build();
}

@GetMapping("/chat")
public String chat(@RequestParam String query) {
    return chatClient.prompt()
        .advisors(
            // RAG Advisor with custom tuning
            QuestionAnswerAdvisor.builder(vectorStore)
                .searchRequest(SearchRequest.builder()
                    .topK(4)          // Fetch 4 chunks
                    .similarityThreshold(0.7) // Only high relevance
                    .build())
                .build()
        )
        .user(query)
        .call()
        .content();
}
}
```